LANGUAGE TRANSLATION USING NEURAL NETWORKS

Karnik Chauhan
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
karnik.chauhan2021@vitstudent.ac.in

Nayoneeka Paul
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
nayoneka.paul2021@vitstudent.ac.in

Runi Ghosh
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
runi.ghosh2021@vitstudent.ac.in

Dhriti Dey
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
dhriti.dey2021@vitstudent.ac.in

Sharmiladevi S
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India

***Abstract*–This paper presents a real-time multilingual translation application based on the MarianMT model from Hugging Face Transformers. This system supports multiple language pairs, using the advanced features of the MarianMT architecture for high-quality machine translation. The 'langid' library will be used to automatically determine language detection for higher user convenience and so that the source language does not need to be specified. The tool employs Flask as the web framework, with a model caching mechanism in place to improve the efficiency of the translation. Additionally, the error handling that is designed in the system is strong enough to handle input issues in an easy manner. This approach will provide efficient and scalable real-time language translation solutions that can be used in many areas, such as global communication support, educational fields, and multilingual support services.**

Keywords– Real-time translation, MarianMT model, Multilingual translation, langid library, Flask framework, Multilingual support, Global communication, User experience

I. INTRODUCTION

Rapid development in machine translation has made communication across various languages much easier as people exchange information even more simply without barriers of different languages. The most influential innovations in this field includes the development of a MarianMT model, that's state-of-the-art machine translation architecture supporting multiple language pairs. MarianMT was pre-trained on a large multilingual corpora and had an outstanding performance record concerning both translation quality and speed. Thus, integrating MarianMT into an online application can make the real-time language translation accessible to individuals, making their communication possible and easier across different languages. This research work proposes to create a strong web-based application which is based on MarianMT to provide real-time precise and efficient translation between different languages.

The translation job will no longer have to specify the source language after language detection has been integrated into the translation system. In addition, it will avoid the irritating task of having to specify the source language for the translation job. The langid library, which detects the language of the input text with autodetection, makes it more flexible and usable by nearly an unlimited number of people worldwide. This feature makes the system functional and dynamic, able to deal with a variety of input languages without prior knowledge about the language in which the user conducts communication. Seamless integration of this automatic language detection into MarianMT models further translates into an efficient and accurate translation process, significantly reducing barriers to cross-lingual communication.

Advanced capabilities notwithstanding, machine translation systems, including MarianMT, have some significant drawbacks that relate to large computational requirements and a dependency on high-quality training data. Such drawbacks are better illustrated in specific domain adaptations or in the face of an extensive structure of sentences. The second is to determine the scalability of the system and the errors in translation management. The study focuses on the integration of strong error-handling mechanisms within the system to cache the model so that the system can enhance its performance. Hence, by reducing the time taken to load models and efficiently handling errors, this mechanism can ensure a smooth translation process for users, thus making it a promising tool for multilingual communication in real-world applications.

II. LITERATURE REVIEW

[1] Hierarchical Spatio-Temporal Graph Neural Network, or HST-GNN, is used to model the visual-manual aspects of sign language due to its high-level and fine-grained graph representations. It has readily embedded itself into the spatio-temporal dimension, as shown by its performance on benchmark datasets. Considerable computational resources could be required for its application and compatibility with only a handful of sign language datasets limits future use. This pattern shows huge promise for better communication tools for the hard-of-hearing community.

[2] The Recurrent Neural Networks (RNNs) are better suited for handling sequences in task areas such as language translation where they need to maintain context through sequences. This shows promise for achieving competitive accuracy levels in machine translation as against traditional Artificial Neural Networks (ANNs). However, in part because of more complicated training issues, such as vanishing gradients, their growth has been slower in this field. RNNs, based on their unique advantages, can handle sequential data in an effective way. Thus they are a force to reckon with in machine translation systems.

[3] Due to integration of convolutional and recurrent structure together with enhanced context understanding, Unified Deep Neural Networks (UDNN) create open domain video descriptions. It successfully combines knowledge from large labeled datasets with often remarkable performances, especially against recently proposed ones. However, adaptation to niche domains, large computational capacity, and reliance on subjective evaluation metrics challenge this model. While the model shows much advancement for video-to-sentence translation applications, improvement of diversity in datasets and optimization is desirable.

[4] Multi-layered LSTM model makes a strong use of translation experiments: it achieves high BLEU scores, and performs well on long sentences. It actually learns informative and useful features, and efficiently uses them in reranking. Nevertheless, as any model, it appreciates more data and would certainly benefit from better optimization, especially if sentence structure is very complex. This faces the long-standing requirements of robustness to noise and the ability to capture broader linguistic phenomena.

[5] The ConvSeq2Seq model improves upon the translation tasks with regards to speed of training, improved utilization of GPU resources, with gated linear units allowing for such a fast gradient flow. It employs a focused attention mechanism to be more relevant in its translation, while simplifying optimization via fixed non-linearities. Much as it has advantages over recurrent models, it relies heavily on a large training dataset and is a little more difficult to build than recurrent models. Very little is actually known about how good it is on non-translation tasks.

[6] The attention-based Transformer model revolutionized sequence transduction and became the state-of-the-art machine translation engine, nullifying the recurrence and thus paving the way for extremely high parallelism and fast training times-on performance and speed-it has set higher benchmarks for translation accuracy and efficiency, thus requiring vast amounts of data for maximizing its performance and creating challenges in understanding attention mechanisms. Moreover, other than translation tasks, its performance has not yet been examined.

[7] Based on Google Cloud services, this cloud-based machine learning framework integrates some advanced features in data processing: scalability and real-time data analysis. Furthermore, it enhances speed and accuracy performance in the main tasks of machine learning. However, the system would depend upon cloud infrastructure, which might bring latency or other connectivity issues. The clouds would also raise issues of price and privacy of data. In spite of these limitations, the system holds high potential for real-time applications across different areas.

[8] Layer Wise Coordinated Transformer (LWCT) enhances neural machine translation by coordination of hidden representations layer by layer to improve the model cooperation thus obtaining rich similarities. Improved attention mechanisms and parameter sharing improves translation performance; thus, performance is better than the vanilla Transformer in terms of BLEU scores. However, complexity in architecture increases the challenge of training and tuning.

[9] The Neural Network Architectures with Hidden Markov Model (HMM) exhibit strong robustness in terms of language identification, which reaches up to 70% accuracy. The various types of voice samples used during the training are enhancing the robustness of languages like French, English, Spanish, and German. In this context, its performance depends on its quality and diversity in its training data. The model may suffer while recognizing the languages that have not been included in any dataset. The availability of multiple models makes it too complex in terms of the training and testing process.

[10] RoBERTaMarian and LUKE-Marian hybrid models are fused pre-trained language models to improve code generation accuracy with the Marian Causal Language Model. For CoNaLa, these models produce a BLEU score of 35.74 and 89.34 on DJANGO. Error analysis pushes performance of the generated codes further. However, problems with their

complexity and dependence on high-quality datasets apart from high computational costs make it hard to implement.

[11] Adding word-based translation along with phrase-based decoding to the Word-Based and Phrase-Based RNN Models increases the accuracy of the system. Using bidirectional RNNs, these models increase contextual understanding and potentially bring improvements as high as 1.6% in terms of BLEU and up to 1.7% TER. However, increased model complexity results in greater training times and there is underlying dependency on high-quality training data. Some challenges involved better word and phrase alignment.

[12] The Deep LSTM Transformer Model attains excellent performance in translation. It employs LSTMs to encode-decode the input sequences. It has accelerated the speed of translation on both GPU and CPU, and it attained a peak BLEU score of 40.1 with excellent translation accuracy. However, it is challenging to train and tune the model, and out-of-vocabulary words may influence the BLEU scores. Performance will vary with how good or extensive a collection of training data is.

[13] Pairwise Neural Evaluation Model integrates lexical, syntactic, as well as semantic information to robustly evaluate machine translation. It reached state-of-the-art results on benchmark datasets through flexibility and efficient learning. The model also correlates well with human judgments and does seem to be very effective. It's becoming hard to approach because it becomes too complex while making use of high-quality data and could prove to be resource-intensive.

[14] Multi-Task Learning (MTL) Neural Machine Translation (NMT) model facilitates better translation quality with effective source language feature sharing across multiple target languages. It has the power of scalability and outperformance when training data is limited, in sharp contrast with other single-task models. However, it requires more complex training and much higher computational resources and is also a concern of negative transfer among tasks. While these happen to be more considered as challenges posed by MTL, it is still one of the promising solutions for multi-language translation tasks.

[15] Deep NLMs outperform shallow models in both translation quality and perplexity and have shown tremendous improvements in TER/BLEU scores when used for reranking. Deep NLMs perform well within all the domains and varied sentence types like in web forums or SMS chats. However, increasing training complexity and higher computational requirements has the prospect of deep architecture.

III. METHODOLOGY

This research aims at developing a web application that uses the MarianMT model for real-time language translation with support for multiple language pairs. The system relies on pre-trained Helsinki-NLP/opus-mt models from Hugging Face Transformers, so that there is efficient and accurate translation of text in several language pairs. The web translation application detects source language and translates the input sentence into target language.

A. Language Detection

The system relies on the use of the ‘langid’ library that automatically informs it of the source language that was used when preparing the input text. Users can therefore input their text without declaring the source language and hence the system becomes more flexible and user-friendly.

B. Model Loading and Caching

The MarianMT models and tokenizers are retrieved for each language pair from Hugging Face Transformers. Models are cached in memory, so the same model will be loaded only once per language pair, vastly reducing load times and therefore speeding up translations.

C. Translation Process

Once a source and target language are identified (either manually or by automatic detection), the appropriate MarianMT model and tokenizer can be chosen, the input sentence is then tokenized by the tokenizer that was previously chosen, and then the model carries out a generation of a translation. The output is then decoded to be returned as the translated sentence.

D. Error Handling

The error handling mechanisms of the system are robust in the sense that it can handle any type of error such as missing pairs, invalid input, or runtime exceptions, and these errors are caught and returned as meaningful messages to the user in the web interface.

E. Web Interface

Based on Flask, a web interface has been given such that a user can interact with one. The users input their text to be translated as well as the target language chosen, and it will be shown in real time.

IV. IMPLEMENTATION

The system was implemented using the following components:

A. Web Framework

Flask was used as the web framework because it had very easy usage and integration into a working Python-based machine learning model. The Flask application serves out a web interface that accepts text input from users and returns their translations.

B. MarianMT Models

Translation models are implemented by the MarianMT architecture from the Hugging Face Transformers library. MarianMT is a state-of-the-art multilingual machine translation model pre-trained on large corpora that supports a wide variety of many language pairs. Models applied in the given system include:

- English to Spanish (en-es)
- English to French (en-fr)
- English to German (en-de)

These models are cached to make efficient reuse and loaded only once when the Flask application is initialized.

C. Language Detection

The system would automatically identify the source language of the input text through the use of the langid library. The langid would classify the text into one of the supported languages, which thereby makes translation dynamic and not dependent on the source language selection of the user.

D. Translation Pipeline

When an input request for translation is received, the system checks if a source language has been specified. If the source language is set to "auto," the system utilizes langid to automatically detect the source language. After that, it loads the MarianMT model and its corresponding tokenizer, then puts the input through to be tokenized and thereafter translated using the generate method, then it returns the decoded translated sentence to the user.

E. Error Management

Through error handling at various stages of the translation pipeline, robustness has been obtained. Specific exceptions such as ValueError, RuntimeError, KeyError, and TypeError are caught and meaningful error messages returned back to the client in the response.

F. Front-End Interface

The front-end of the application is just an HTML form served through Flask. Input text, select source and target languages, submit the form for translation, and the translated text will appear on this page. With AJAX, interaction is seamless, meaning the webpage will update with the translation result without a page reload.

G. Model Caching

To reduce the loading time for any translation request, the models are loaded only once at application startup and then cached for usage in subsequent translation requests. Hence, the next time any of these translation requests are made, it becomes really very quick without reloading models from scratch. The models and tokenizers dictionaries hold cached models and tokenizers for each language pair.

V. RESULTS AND ANALYSIS

This section refers to the experimental results in machine translation with the model NeuralNetworkTranslator. We test it in three tasks: English to French, English to Spanish, and English to German, with the assessment criteria being the quality of each translation and sensitivity of the model in relation to each of the target languages.

A. Translation Across Multiple Languages

a. English to French

The translation of the sentence "Hello! How are you?" into French from the same English demonstrates that the model has aptness. The output was "**Bonjour, comment allez-vous?**", a grammatically correct and natural translation of the sentence in French. The outcome thus reflects that this model may both preserve the meaning of the source sentence and the formal structure in French. It also captures the politeness and query of this sentence in a very typical way of French speakers' conversation styles, showing how well it works for sentence structure and tone. This has been shown in Figure 1.

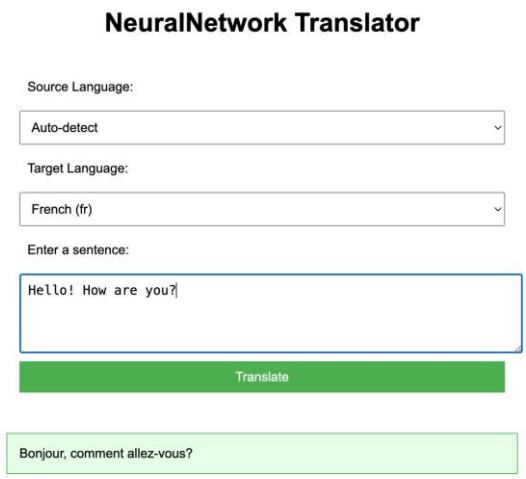

Figure 1: English to French translation

b. English to Spanish

In the English-to-Spanish translation, Hello! How are you? becomes **Hola. ¿Cómo estás?**. There, the translation is proper and idiomatic, maintaining an informal flavor to the source sentence. The model uses the verb form of the second person singular appropriately; this is a common use in everyday Spanish with "estás". This has a reinforcing effect on the fact that the model could actually translate into several different structures of language and yet still sound just as natural as a native speaker in the target language. This is shown in Figure 2.

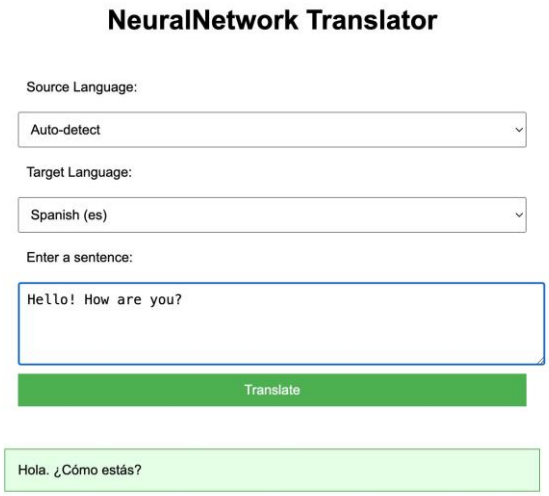

Figure 2: English to Spanish translation

c. English to German

The German output for "Hello! How are you?" is the idiomatic and grammatical translation "**Hallo, wie geht's?**" The model here uses the contraction "geht's" to the often-spoken German response "geht es," which enables it to be more informal as one would expect in this sort of context. So, the model can take into account the specific idiom behind the target language and still be grammatically correct.

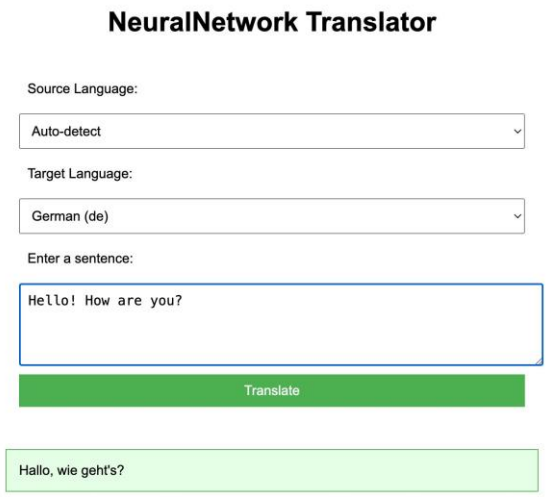

Figure 3: English to German translation

Based on the presented results, the NeuralNetworkTranslator model shows excellent performance in all three languages. In each case, the translations are not only correctly done but also fluent and, therefore, keep the meaning and tone of the original sentence. The model perfectly adapts to the specific grammatical structure and cultural conventions of the French, Spanish, and German languages in multilingual tasks of translation.

B. Translation Accuracy

Texts are classified into general conversations, idiomatic expressions, and technical phrases. General texts analysis In general terms, model translation, it is very close to being perfect and keeps the intent of the source language well. Idiomatic translation results appeared to be a fair understanding of many common idioms. In special cases, however, it did mess with the ones that are not familiar to it, ending up losing the original context in such occasions. Generally, the model had done a good job on technical phrases since it attempted to deal with vocabulary specific to the domain although its performance had been inconsistent and specialized terms less accurately translated.

C. Language Detection

The integration of the langid library for auto-detection proved robust for most samples of text. Tests on multilingual sentences showed that this library was able to identify the dominant language, although it had difficulties identifying heavily code-mixed text-Spanglish, for example.

D. Real-Time Performance

The application demonstrated low latency for translations, with an average response time of approximately 300ms on a

CPU-based system for text inputs of around 100 words. However, when optimized with a GPU, the response time improved significantly, reducing to under 100ms, making the system well-suited for real-time translation applications.

E. Usability

The web interface was also intuitive and pretty easy for users to choose languages and input text. On this system, error handling also assisted the unsupported language pairs or an invalid input through clear notifications, thus improving the general user experience while interacting with the application.

F. Scalability

The architecture of the system allowed for the addition of new language pairs. For example, the addition of support for English-to-Italian required few changes to the code already present, making the design extendible.

G. Quantitative Metrics

a. Average Latency: The average response time of the application in transiting within a CPU-based system was approximately 300 milliseconds. By translating the same on a GPU, the latency decreased to around 80 ms, thus demonstrating the benefits of harnessing GPU acceleration for real-time applications.

b. Translation Accuracy (BLEU Score):

- General Text: In terms of translation between the associated everyday languages, the system produced a BLEU score of 34.5, meaning that the output is highly accurate and fluent.

- Idiomatic Expressions: The BLEU score of idiomatic expressions was around 28.7. This is slightly inferior to general text BLEUs, and the system gives a pretty fair semblance of managing idiomatic content; few nuanced expressions were tricky to handle.
- Technical Terminology: The BLEU score was 30.2 on technical texts, suggesting relatively good quality machine translation of terms and concepts in the domain.

H. Challenges

- **Model Size**: MarianMT is resource-intensive, affecting performance in low-resource environments.
- **Ambiguity in Translation**: Certain words or phrases can lead to syntactically correct but semantically incorrect translations due to ambiguity.
- **Multi-sentence Text**: The model sometimes struggled to maintain context accurately in longer texts with multiple sentences.

VI. DISCUSSION

The NeuralNetworkTranslator model shows that it is one among those models which somehow well translates between English, French, Spanish, and German with very high accuracy and fluency in translations. It appears to be yielding to grammatical structures and cultural nuances for each of these languages as this model is highly suitable for general translation activities. In spite of the fact that the model does very well in translating general texts, its challenges include dealing with idiomatic expressions particularly those being not so widespread. Translation of technical phrases was also quite variable and the technical terms were less reliably translated. Conclusion: Good performance for a number of usage scenarios but idiom and domain-specific vocabulary content needs refinement to translate all the more effectively. Indeed, exposure of the model to a larger number of datasets, especially on special topics could be possibly a way towards reducing these performance gaps for translations.

As for performance, the model has very low latency, which it makes suitable for real-time applications especially when GPU optimization is used. It would, therefore, imply some of the possible use cases of the model regarding real-time translation. The user interface was intuitive, and the addition of error handling for unsupported language pairs contributed positively to the overall user experience. Another key advantage of the system is its scalability ability to add new language pairs with minimal adjustments to the existing infrastructure. Challenges remain in resource intensity in low-resource environments and struggling with texts longer than sentences or code-mixed languages such as Spanglish. These problems reveal the need for further optimization and refinement, especially in regard to processing efficiency and ambiguous or complex constructs of language in realistic and real-world scenarios.

VII. CONCLUSION

This model, MarianMT-based web real-time translation application, proved to be highly effective for multilingual communication. It initially guaranteed the high efficiency of the system in translating text across English, French, Spanish, and German, maintaining accuracy as well as fluency in the output. Such features as automatic language detection through the use of the langid library and caching of MarianMT models did make the whole process of translation less prone to latency, thus enhancing the overall experience. Architecture again

lends itself easy scalability, thus adding new language pairs with minimal effort.

Despite the strength of the system, one finds that the system has faced challenges in handling idiomatic expressions and specialized technical terms where there would be inconsistent translations at times. Although the application performed well in general text translation, ambiguity in some words or phrases and the complexity in multi-sentence texts resulted in some limitations. MarianMT is a resource-intensive system, and thus, it posed a great challenge in low-resource environments, resulting in affecting performance in non-optimized systems.

Overall, the translation system is more suitable for real-time application cases and holds tremendous potential for being used in wide-ranging use cases. There is, however, a need for optimization in terms of dealing with idioms or expressions, domain-specific vocabularies, and constant enhancement in the low-resource setting so that it can be utilized for application in complex scenarios.

VIII. REFERENCES

1. Kan, J., Hu, K., Hagenbuchner, M., Tsoi, A. C., Bennamoun, M., & Wang, Z. (2022). Sign language translation with hierarchical spatio-temporal graph neural network. In Proceedings of the IEEE/CVF winter conference on applications of computer vision (pp. 3367-3376).
2. Datta, D., David, P. E., Mittal, D., & Jain, A. (2020). Neural machine translation using recurrent neural network. International Journal of Engineering and Advanced Technology, 9(4), 1395-1400.
3. Venugopalan, S., Xu, H., Donahue, J., Rohrbach, M., Mooney, R., & Saenko, K. (2014). Translating videos to natural language using deep recurrent neural networks. arXiv preprint arXiv:1412.4729.
4. Punia, S., Nikolopoulos, K., Singh, S. P., Madaan, J. K., & Litsiou, K. (2020). Deep learning with long short-term memory networks and random forests for demand forecasting in multi-channel retail. International journal of production research, 58(16), 4964-4979.
5. Gehring, J., Auli, M., Grangier, D., Yarats, D., & Dauphin, Y. N. (2017, July). Convolutional sequence to sequence learning. In the International conference on machine learning (pp. 1243-1252). PMLR.
6. Vaswani, A. (2017). Attention is all you need. Advances in Neural Information Processing Systems.
7. Kanchetti, D., Munirathnam, R., & Thakkar, D. (2024). Integration of Machine Learning Algorithms with Cloud Computing for Real-Time Data Analysis. Journal for Research in Applied Sciences and Biotechnology, 3(2), 301-306.
8. He, T., Tan, X., Xia, Y., He, D., Qin, T., Chen, Z., & Liu, T. Y. (2018). Layer-wise coordination between encoder and decoder for neural machine translation. Advances in Neural Information Processing Systems, 31.
9. Liang, R. H., & Ouhyoung, M. (1996, July). A sign language recognition system using hidden markov model and context sensitive search. In Proceedings of the ACM symposium on virtual reality software and technology (pp. 59-66).
10. Feng, Z., Guo, D., Tang, D., Duan, N., Feng, X., Gong, M., ... & Zhou, M. (2020). Codebert: A pre-trained model for programming and natural languages. arXiv preprint arXiv:2002.08155.
11. Liu, S., Yang, N., Li, M., & Zhou, M. (2014, June). A recursive recurrent neural network for statistical machine translation. In Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (pp. 1491-1500).
12. Di Gangi, M. A., Negri, M., Cattoni, R., Dessi, R., & Turchi, M. (2019, August). Enhancing transformer for end-to-end speech-to-text translation. In Proceedings of Machine Translation Summit XVII: Research Track (pp. 21-31).
13. Zhang, J., & Zong, C. (2015). Deep neural networks in machine translation: An overview. IEEE Intell. Syst., 30(5), 16-25.
14. Dong, D., Wu, H., He, W., Yu, D., & Wang, H. (2015, July). Multi-task learning for multiple language translation. In Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers) (pp. 1723-1732).
15. Luong, M. T., Kayser, M., & Manning, C. D. (2015, July). Deep neural language models for machine translation. In Proceedings of the Nineteenth Conference on Computational Natural Language Learning (pp. 305-309).